


ClaudeClaw V3 Visual Guide

The diagrams from the V3 video, compiled with notes

Read [DISCLAIMER.md](#) before relying on anything in this kit.

01 The Wrapper


The big picture in one image



- Your local computer is the engine. Claude Code runs as the brain.
- The Hive Mind is the wrapper around Claude Code that orchestrates agents, memory, and skills.
- Telegram is the steering wheel — your portable interface that exposes the wrapper anywhere you have your phone.
- Slack and Discord work too. Telegram is the most reliable for on-the-go.
- The brain is replaceable. Tomorrow when a better brain exists, you swap it. The wrapper stays.

02 The Meta Agent Example

How a CLI plugs into the ecosystem




- Install the CLI globally on your machine (e.g., `npm install -g meta-ads-cli`).
- Add a skill folder pointing at the CLI. Document what it does in `SKILL.md`.
- Every agent in your Hive Mind can now use the CLI through their tool allowlist.
- Schedule it via cron to fire on a timer. Now you get a Facebook ads brief at 7:30 AM every weekday.
- The pattern works for any CLI. Add it once, every agent inherits it.

03 What an AI OS Actually Is

The modular brain frame


WHAT AN AI OS ACTUALLY IS.



- An AI OS is the layer between your ecosystem and the brain.
- Bottom layer: your folders, files, scheduled jobs, integrations. Already exists.
- Middle layer: the orchestration. Skills, agents, memory, the war room. This is what you build.
- Top layer: the brain. Replaceable. Today it's Claude Code. Tomorrow could be Codex. Could be both.
- AI OS is the layer. The brain is replaceable.

04 Mission Control

Kanban + auto-assign




- Tasks live in three columns: Queued, Running, Done.
- When a task arrives without a target, a Gemini classifier reads the task and routes it to the right agent.
- The classifier sees the task description plus the agents' descriptions. It picks the best fit.
- Right agent picks up the right work. You don't have to think about routing.
- The kill switch
`MISSION_AUTO_ASSIGN_ENABLED` gates this if you want manual routing only.

05 Scheduler

Cron, with a translation layer

THE SCHEDULER.




Behind the scenes is just cron. **You never see it.**

- Behind the scenes: standard cron expressions.
- In front of you: friendly English (every morning at 6 AM, weekdays at 6:30 PM).
- The dashboard translates one to the other so you don't have to learn cron syntax.
- Power users can edit the raw cron in an Advanced toggle.
- The kill switch `SCHEDULER_ENABLED` stops all scheduled task execution without restarting anything.

06 Agents Tab + Suggestions

Where you manage the team


AGENTS TAB.



- Every agent is a folder with two files: agent.yaml (model + tools) and CLAUDE.md (persona).
- Edit either file directly in the dashboard. No terminal needed.
- The Suggestions feature watches your usage. When one agent is overloaded, it suggests a split.
- Example: "Your Comms agent is handling email, Slack, WhatsApp, and YouTube comments. Consider an Email Manager agent."
- Suggestions are nudges, not auto-execution. You decide.

07 Chat


All channels into one view



- Telegram, Slack, Discord, web — every channel funnels into one unified chat history.
- You can scroll your whole history without remembering which app you were in.
- All channels respect the same agent routing, memory, and tool allowlists.
- The bridge is where the message comes in. The agent and memory don't care.

08 Memories

Three-layer hybrid recall




Three layers stacked. Old context fades. **Important context survives.**

- Layer 1: FTS5 keyword search. Fast. Misses semantic similarity.
- Layer 2: Embedding-based semantic search. Catches meaning even if keywords don't match.
- Layer 3: Saliency scoring. Memories that have been useful in past responses get a boost.
- The three layers are merged at query time. The agent sees a ranked list of relevant memories.
- Old unused memories decay. Pinned memories never decay.

09 Hive Mind Views

Three perspectives on the same data

THE HIVE MIND.




Same data. Three ways to see it.

- List view: chronological log of what every agent has done. Simplest, always works.
- 2D graph: Obsidian-style force-directed graph of tasks and shared knowledge.
- 3D brain: cinematic visualization. Brain regions light up as agents fire.
- All three views read the same underlying `hive_mind_log` table.
- If the list view is operational, you can layer the 2D and 3D views on top later.

10 Audit


Append-only trail of every state change



- Every tool call by every agent. Every kill switch flip. Every dashboard config change.
- Append-only. 90-day retention. Pinned rows survive prune.
- Correlation IDs group related events from one user-initiated turn.
- When something happens, you can always trace it back. No silent state changes.
- Pair with the kill switches and the exfiltration guard for layered safety.

11 The War Room

Multiple agents, one prompt




- Five agents seated around a virtual table: Main, Comms, Content, Ops, Research.
- `/standup`: every agent gives a brief morning report. What they wrapped, what's queued, what's blocked.
- `/discuss`: every agent weighs in on a question with their unique angle. Main consolidates.
- Each agent runs in isolation. They don't see each other's responses. The consolidator runs last and sees everything.
- Voice mode also exists. Same patterns, different bridge.

12 An OS Is a Data Engineering Exercise

The thesis in one image

AN OS IS A DATA ENGINEERING EXERCISE.



Right files. Right place. Right time.
That's the whole game.

- Left: most people's setup. Files everywhere. The LLM has nothing to grab onto.
- Right: the Hive Mind. Three folders and a database. The LLM knows where to look.
- Right files. Right place. Right time. That's the whole game.
- Everything else in this kit is decoration on top of this idea.
- If you understood this image, you understood the architecture.

13 Project vs Global


The cleanliness rule



- Two levels: global (defaults available everywhere) and project (overrides for one project).
- The rule: project always wins on conflict.
- Set sensible defaults at the global level. Override at the project level only when you need to.
- Same idea applies to skills, agents, and CLAUDE.md.
- A clean house makes everything else achievable.

14 Add Your First Agent

Five steps to grow your team




- Step 1: Copy the `_template` folder.
- Step 2: Rename to your agent's name.
- Step 3: Edit `agent.yaml` — set the model, set the tools allowlist.
- Step 4: Edit `CLAUDE.md` — write a one-paragraph persona.
- Step 5: Get a Telegram bot token from `@BotFather`. Save it in `.env`.
- Refresh the dashboard. Your new agent is live.

15 Where Everyone Starts

Step 1 of the journey

WHERE EVERYONE STARTS




Files everywhere. **No system.**

- PDFs piled on screenshots piled on text files piled on random folders. No system.
- If you've ever opened your laptop and felt that little wave of dread, this is why.
- The reason you don't have a working AI OS yet is not because you haven't read enough about AIOS.
- It's because the file system underneath any AI OS would be feeding off this chaos.
- Step 1 is admitting this is the starting line.

16 Group the Files

Step 2 of the journey

GROUP THE FILES.




- Six folders: Documents, Images, Videos, Code, Notes, Assets.
- Sort what's already on disk into one of these buckets.
- Don't worry about subfolders. Don't migrate cloud storage. Don't rename anything.
- Same files. Just now they have a home.
- This step alone gives the next four steps somewhere to live.

17 Untangle the Tools

Step 3 of the journey


UNTANGLE THE TOOLS.



- Three categories of tooling: skills, rule files, CLI tools.
- Probably tangled together right now. Untangle them.
- Global skills and global CLAUDE.md live in ~/.claude/. Project overrides live in the project's .claude/.
- Sounds boring. It's the unlock.
- Step 3 is where most people stop. Don't.

18 Build the Hierarchy


Step 4 of the journey



- Two levels: global and project. Project always wins on conflict.
- Set sensible defaults at the global level. Override at the project level only when needed.
- This is the rule that keeps your Hive Mind organized as it grows.
- Same pattern for skills, agents, and CLAUDE.md.
- A clean house makes everything else achievable.

19 Layer on the Bridge

Step 5 of the journey



- Bottom layer: your organized ecosystem (steps 1-4).
- Middle layer: ClaudeClaw / your Hive Mind orchestration. Agents, skills, memory, scheduler.
- Top layer: the bridge. Telegram is the recommended portable interface. Slack and Discord work too.
- As long as your computer is on at home, the bridge keeps the whole thing running anywhere you are.
- That's the full architecture. Five steps. Anyone can do this in a weekend if they care enough to start.