

Use more models. Keep your workspace.

Codex + Ollama Setup Kit

Keep a capable model available for difficult work, and give a cheaper cloud or local model one well-defined job. This guide helps you connect it, prove that it can use tools, and decide whether it is useful on your machine.

Your first finish line

A fresh Codex task uses your chosen model to calculate 137×249 , write 34113 into a real file, and read it back. For a local route, you also verify the model running in Ollama.

Three things from the video

A copy-ready setup prompt. A practical guide. The workflow and economics website. Everything is linked from START-HERE.md in the ZIP.

Choose a route

Local: use a model downloaded to your computer. Choose this when it fits your hardware and your task can tolerate its speed and quality.

Cloud through Ollama: use a hosted model. Choose this when you want another model without loading its weights on your computer. Check its allowance and cost.

Your existing OpenAI model: keep it available for planning, difficult reasoning and review. The kit does not replace your subscription or reset an allowance.

Version 1.0 | 16 September 2026 | Companion to video 206

Connect the right version of the app.

Start with your existing working setup. Save any open work. Install Ollama from its official download page if needed, then use the guided setup prompt to inspect your machine before choosing a model.

1. Check the desktop connection

The current Ollama documentation calls this ChatGPT Desktop with Codex mode. It specifies Ollama v0.34.0 or newer: open Ollama Apps, enable ChatGPT (Desktop), and follow its setup prompts. Restart the desktop app when asked. [1]

2. Choose the model you will actually use

In current documentation, Ollama Settings > Apps > ChatGPT lets you select compatible local or cloud models. Save your choice, complete any requested restart, then select the intended model in a fresh Codex task. [1]

The recording may show Add to Codex or Add to ChatGPT. Labels and available models depend on your installed versions. If your screens differ, follow the current official desktop guide rather than inventing settings.

3. Keep the terminal route separate

Ollama also documents **ollama launch codex** for Codex CLI. That is a terminal integration; a successful CLI test does not prove your desktop picker has updated. [2]

4. Know how to disconnect

The current desktop guide says to turn off ChatGPT in Ollama Apps and restart when prompted to restore the normal connection. Record any other configuration changes separately. [1]

Official desktop setup: docs.ollama.com/integrations/chatgpt

Download: ollama.com/download

Give Codex a clear setup job.

For your own machine, copy the full text from **02-GUIDED-SETUP-PROMPT.txt**. It asks Codex to inspect, propose, connect and verify in that order. Review the proposed model and changes before proceeding.

The original website prompt

I've installed Ollama. Inspect my computer's memory, GPU, free disk space and installed models. Choose a tool-capable local model that fits comfortably while leaving room for my other apps. Reuse a suitable installed model, otherwise explain the download size and install a suitable official model. Connect it to Codex using the supported Ollama integration, preserve my current cloud models and default, and back up any configuration you change. Verify local inference, then use Codex to edit a test file and read it back. Tell me what worked and whether I need to restart the app. Don't claim success from a dropdown entry alone.

What to expect back

A hardware summary, one model recommendation plus a smaller fallback, local/cloud route, download size, expected context and memory needs, configuration changes and restart steps. A model recommendation without these details is not ready to act on.

Why the guided version adds a review step

A large download can consume disk space, a cloud route can incur charges, and configuration changes can affect your default model. The guided prompt asks for a concrete proposal first and preserves a route back to your working setup.

The original prompt is also in 01-SETUP-PROMPT-FROM-SITE.txt. Neither prompt has been verified as a universal fresh-install procedure.

Prove the model can do the work.

Open a disposable project or choose an unused test folder in a project. Select the model explicitly. Copy **03-VERIFY-YOUR-MODEL.txt** into a fresh task.

The smallest useful tool test

Calculate 137×249 with a tool. Write the result to result.txt. Read the file back with a separate tool call. Expected contents: 34113.

Check the file yourself and inspect the tool output. A correct answer typed in chat is not evidence that the model wrote or read a file. Do not let the task silently switch to a different model.

If you chose a local model

While the request is running, inspect **ollama ps**. Check the exact model and its processor split; the context column can help confirm allocation. CPU/GPU placement affects speed, and a larger context needs more memory. [3, 4]

Record your result

Model and tag: _____

Local or cloud: _____

Test file and contents: _____

Tool evidence: _____

Elapsed time and retries: _____

Then test one real task

Give it a short source file to summarize, a small edit with an exact expected result, or a shortlist from provided facts. Compare the finished output with your acceptance checks before using it more broadly.

Passing this test proves a narrow capability in your current setup. It does not prove universal tool, browser, skill or agent compatibility.

Pass small jobs between models.

The video demonstrates cloud research, a local shortlist, and a final synthesis. Reuse the division of work on a task you can check. The exact models are examples, not a required roster.

Plan

Ask your stronger model to define the outcome, allowed files, required inputs and acceptance checks. Put the plan in a project file so another task can read it.

Do one bounded piece

Give the chosen cloud or local model one job: extract fields, compare provided options, rewrite a small section, or implement a narrow change. Keep its context relevant and set a retry limit.

Review the actual result

Have the reviewer inspect the files, calculations or tests. A confident summary from the worker is not enough. Return one precise correction if the result misses a check.

Use 04-TASK-HANDOFF-TEMPLATE.md. Fill in outcome, inputs, allowed changes, acceptance checks and budget before switching tasks.

A practical example

Task 1: research three hardware options with source links and dates. Task 2: shortlist those options against your real memory, budget and portability constraints. Task 3: check the evidence and explain the tradeoff.

Use separate chats or manual model selection if your app cannot create model-specific tasks. Do not assume that changing a model preserves every detail of prior context.

Compare the cost of a finished task.

Open the companion website to revisit the diagrams and economics shown in the video. Treat its dated price bars as a starting point. Check provider prices and your own account before spending.

[Open the companion website](#)

Three different cost buckets

Your OpenAI plan: keep track of its included usage or your separate API bill.

Ollama Cloud: hosted requests draw on its own allowance or purchased credits.

Local inference: no cloud-model token charge for the local computation, but hardware, electricity and your time still have a cost.

Current Ollama plan snapshot

Checked 16 September 2026: Free includes starter usage; Pro is \$20/month with \$60 of usage credits; Max is \$100/month with \$300 of usage credits. These are allowances, not a promise that every model is free or every task costs the same. Recheck ollama.com/pricing. [5]

Use the right denominator

Compare cost per accepted result: model charges + paid tools + retries + your review time. A cheaper token rate can still produce a slower or more expensive finished task.

For uncached usage, a rough model-only estimate is $(\text{input tokens} / 1,000,000 \times \text{input rate}) + (\text{output tokens} / 1,000,000 \times \text{output rate})$. Add other billable usage separately. Do not translate a rate ratio into a guaranteed savings percentage.

Local model inference does not make connected tools local. Web search, external APIs and other integrations may still use network services.

When something does not work.

The model is missing from the picker

Confirm the intended model is selected in Ollama Apps, save, and fully quit/reopen the desktop app when requested. Verify app versions against the current desktop guide. Test in a new task. [1]

It talks but does not use tools

Check that the selected model supports tools in this integration. Retry only the small verification test. If it keeps failing, use a compatible alternative; do not call a conversational response a passing result.

The computer slows down or runs out of memory

Reduce the workload or choose a smaller model that fits at the intended context. Model download size is not total runtime memory. Check the actual allocation and CPU/GPU split. [3, 4]

The request went to the wrong place

Check the selected model and route before retrying. A locally installed Ollama app can call a cloud model. Report route uncertainty rather than assuming a local result.

A model switch loses useful context

Save the task plan, sources and acceptance checks in a file. Start the next task with that file instead of relying on a long chat to transfer perfectly.

You need your previous setup back

Use the supported desktop disconnect flow and any specific backups made during setup. For a CLI launch profile, use its own documented restore procedure; do not apply CLI instructions blindly to desktop. [1, 2]

Stop after two attempts that fail for the same reason. Capture the error, app versions and exact model, then change one relevant condition.

Your resource checklist.

Before you call the setup complete

- ☐ You know which model and provider each task uses.
- ☐ The test file exists and contains 34113.
- ☐ Local execution is checked when you choose a local route.
- ☐ One small real task meets its acceptance checks.
- ☐ You know how to return to your previous setup.

Official references

[\[1\] Desktop integration](#)

[\[2\] Codex CLI integration](#)

[\[3\] Context length](#)

[\[4\] Ollama FAQ](#)

[\[5\] Current pricing](#)

The website shown in the video

[More models. One workspace.](#)

The site supplies visual explanations and a copyable prompt. This PDF supplies the step-by-step viewer guide. The ZIP includes plain-text prompts so you can copy them without PDF formatting.

Keep building

[Build, Learn, Earn with AI](#)

Mark Kashef | 16 September 2026

Source check is dated. App behavior, model availability and prices can change. This is a practical companion, not a performance guarantee.