


Why should you care?



How does this process work (simply)?

Document is broken into pieces (chunking)



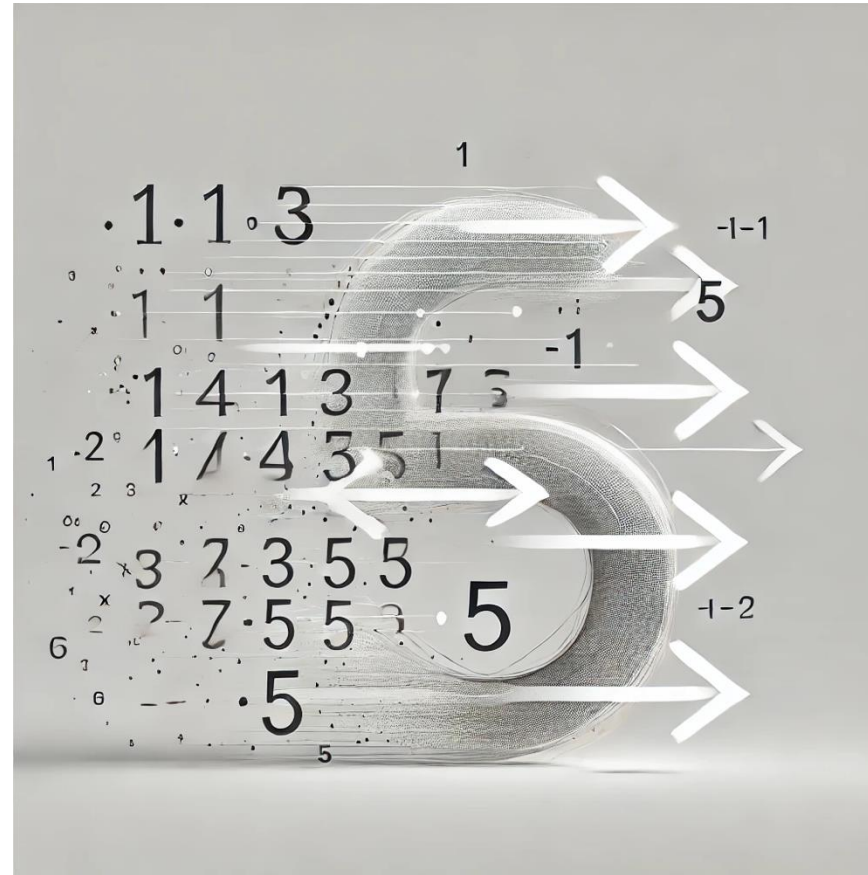
How does this process work (simply)?

These pieces are converted into numbers (**vectorizing**)



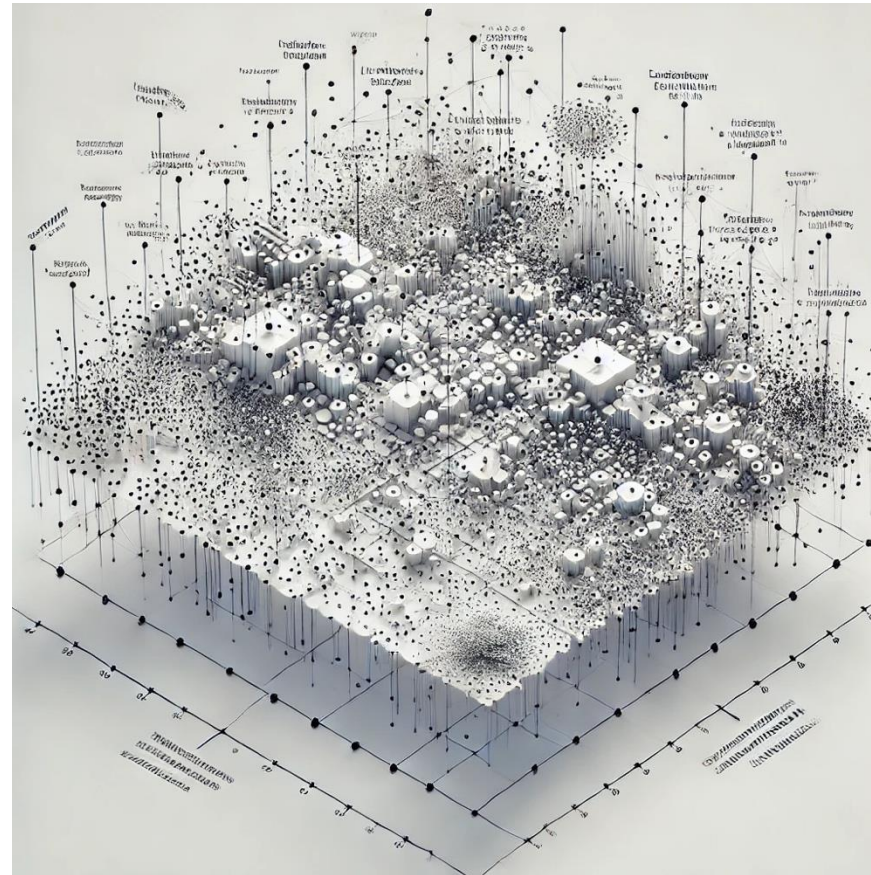
How does this process work (simply)?

The numbers are then converted into a special 'language'



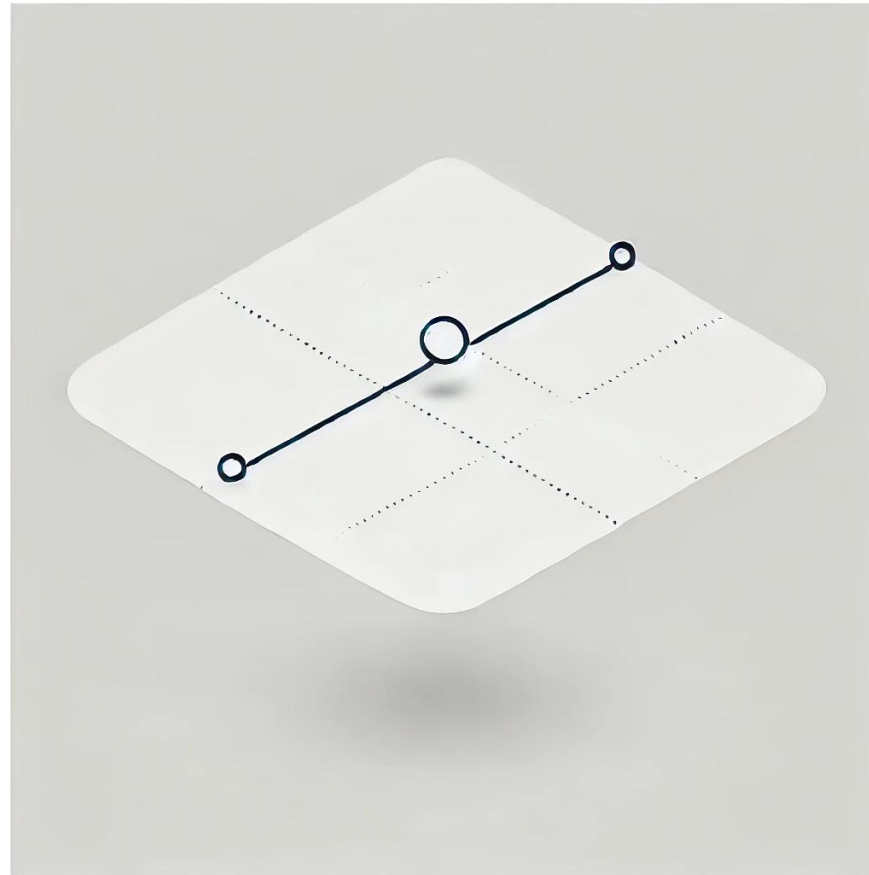
How does this process work (simply)?

The vectors are thrown into a 'map' (**embedding**)



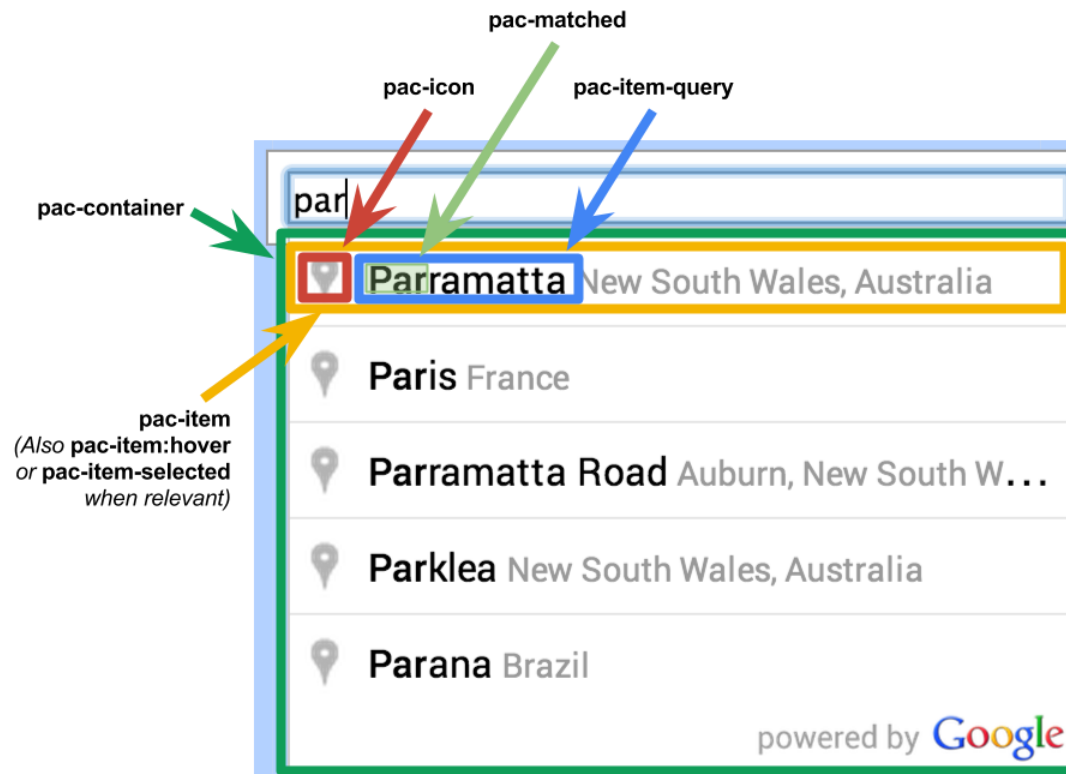
How does this process work (simply)?

To find answers to questions, you pick the closest matches to your 'destination'



How does this process work (simply)?

Think of searching for a destination on Google Maps





YOU: Mark, I repeat, why should I care?

**ME: Because most tools, including GPT,
are BAD at this**



Why are they bad?

They do a quick scan rather than a deep dive.



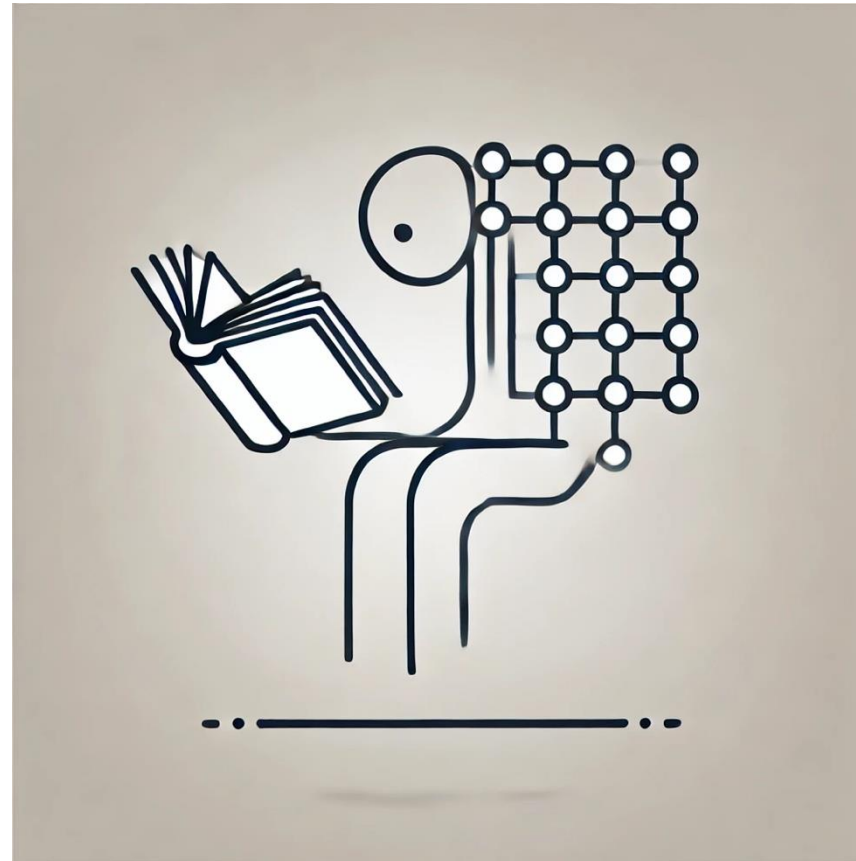
Why are they bad?

Even worse, the LLMs ‘consult’ whatever is found, rather than use it as the ‘ground truth’



How Pinecone Assistant Helps

The system 'talks' to the document before the LLM is ever asked to respond.



**The result is either response from your document,
or ‘I Don’t Know’ (my favourite response)**

Accurate Knowledge Base GPT

