

Define done. Then let it work.

Give your agent an output to produce and a way to prove it finished. Fill this out before handing over the task.

01

The result

What should exist when the work is done? Name the output, who it serves and where to save it.

02

Three passing checks

For each check, name the input, the expected result and the evidence to return. Replace words like better, clean and fast with something you can inspect.

03

Allowed changes

List the files or systems the agent may change. Name what stays untouched and which actions need your approval.

04

When to stop

Set a time or spending limit that your tool can enforce. Stop on missing access, an unclear requirement or the same failed check twice. Return the blocker and work so far.

05

The handoff

Ask for the finished output, a pass/fail/not-run result for every check, the evidence and anything still unresolved. You decide whether to accept it.

Copy the task brief from the included text file. Leave unknown details blank and ask the agent to resolve them with you before it begins.

Written instructions still need account permissions, tool limits and your review. A prompt cannot guarantee that an agent stops or stays within budget.

From vague to verifiable

Instead of “make the page better,” define one repair and three passing checks. This example is a specification you can adapt.

Task and output

Fix the contact form in the local preview. Return a reviewable change and screenshots. Keep the existing copy and design.

Check 1 / Empty email

Submit with the email field empty. Show an email-required message and send no request. Return a screenshot and request log.

Check 2 / Invalid email

Submit abc as the email. Show an invalid-email message and send no request. Return a screenshot and request log.

Check 3 / Valid submission

Submit a valid test address with the test receiver returning success. Exactly one request reaches it and a confirmation appears. Return the request log and screenshot.

Scope and stop rule

Change only the form and its validation. Use synthetic data and the test receiver. Ask before adding a dependency. Stop if the receiver is unavailable or the same check fails twice. Set any time or cost cap in the tool.

Final review

Return the change and each check with pass, fail or not run, plus the evidence. List unresolved issues. Get my approval before publishing to production.

Illustrative example. No execution or passing test result is claimed. Adjust the checks and permissions to your own task before using it.