



The Effort Decoder

The end-to-end guide to the most misunderstood setting in AI. What the effort dial actually does, what it really costs, and a workflow for picking the right level in seconds, on Claude, OpenAI, Grok, or Gemini. Plus the exact control prompts to test it all yourself.



Everyone piles at the expensive end



WHAT EFFORT ACTUALLY IS

A budget for how long the model is allowed to think before it answers. Same model, same knowledge, same intelligence at every rung. Anthropic's own docs describe it as how eager the model is about spending tokens.

THE ADDICTION PATTERN

People start at extra high or max, it works, and they never come back down. Not because lower failed. Because they never risked finding out.


THE INSURANCE PSYCHOLOGY

"I'd rather overkill an easy task than risk a wrong answer." Fair. But you're paying insurance premiums on tasks that were never going to crash.

THE FIX IN ONE LINE

Start low. Climb only when the output tells you to, never on vibes. The rest of this guide is the evidence for why that works.

The top of the dial can hurt accuracy, not just cost



THE ANALOGY

Study too much for a multiple-choice exam and you can justify every option on the sheet. You had it. It was B. And you talked yourself out of it.

THEIR OWN DOCS AGREE

Anthropic, verbatim: max "can lead to overthinking." The vendors that sell the tokens are telling you the top rung can backfire on simple tasks.


WHEN HIGH EFFORT EARNS ITS KEEP

Genuine ambiguity, deep multi-step architecture, and long unattended agentic runs. Anthropic pins extra high to agentic tasks over 30 minutes. Not your emails.

TWO DIFFERENT DAMAGES

The exam is the accuracy damage, overthinking a known answer. The committee that won't adjourn is the cost damage. Two separate reasons not to max by default.

You maxed the jar, not the machine



MODEL VS HARNESS

The model is a brain in a jar. It can think and nothing else. The tools bolted around it, read files, run code, search, edit, that's the harness. Capability mostly lives there.

GOOGLE'S RULE OF THUMB

Google's vibe-coding whitepaper frames an agent as model plus harness, roughly 10 percent model, 90 percent harness. A rule of thumb, not a lab measurement, but the direction has receipts.


THE RECEIPTS

A Terminal-Bench team went from outside the top 30 to the top 5 changing only the harness. A peer-reviewed Princeton study roughly tripled a fixed model's solve rate by redesigning nothing but the harness.

WHY THIS REFRAMES THE DIAL

Cranking effort pays to make the brain sit in the jar longer. It upgrades none of the hands and legs. That's why maxing the dial so often buys nothing.

Four names, four heights, verified July 2026



FOUR NAMES FOR ONE KNOB

Anthropic says effort. OpenAI says reasoning effort. xAI says reasoning effort too. Google says thinking level. Same dial, four labels on the dashboard.

NO TWO LADDERS MATCH

OpenAI runs six rungs (none to max), Claude five, Gemini 3.x four, Grok three. The same word sits at different heights: high is the top of Grok's ladder and the middle of Claude's.


THE DEFAULTS ARE SNEAKY

Leave the dial alone and Claude and Grok run at high, OpenAI at medium, and Gemini decides for itself per request. The do-nothing habit gets you a different bill per vendor.

EVERY NUMBER CHECKED

Every rung on this chart comes from the vendors' own docs, and the disputed ones were verified live against the APIs in July 2026. This is the current state, not folklore.

The center is the mental model, the edges are the traps



ALL FOUR AGREE

Effort is a thinking-token budget. Thinking is billed as output tokens, the expensive kind. More effort means more tokens and more waiting, not a smarter model. And models adapt within a level.

HINT VS HARD BUDGET

Anthropic is explicit that effort is “a behavioral signal, not a strict token budget.” Gemini's legacy path was a hard token number. Same dial name, different contract.


CAN YOU TURN IT OFF?

OpenAI, yes (none). Grok, no. Claude and Google, depends on the model, and on several flagships thinking cannot be disabled at all.

CAN YOU SEE WHAT YOU PAID FOR?

Grok shows the raw trace. OpenAI gives a count and summaries. Google gives summaries. Anthropic hands you signed blocks, sometimes with the content encrypted. You always pay either way.

The consumer surfaces hide, rename, and move the setting



CHATGPT

Extra High is xhigh with a nicer jacket. The Instant option is powered by last generation's model. Max is hidden in settings until you enable it. Auto-switching can quietly change effort mid-conversation.

CODEx

The Codex app displays low as Light and xhigh as Extra High, and it lists Ultra on the effort menu, even though the API itself rejects ultra as an effort value. The vendor's own tools disagree.


GROK

Grok's app shows no effort dial at all. Four modes instead: Fast, Auto, Expert, Heavy. Heavy is literally subtitled Team of Experts, a multi-agent fan-out sold as a menu item.

THE OTHER TWO CORRECTIONS

Effort is not temperature, deliberation and randomness are different knobs, and temperature is locked on the newest models anyway. And no rung adds knowledge: low and max share the same training cutoff.

What low, medium, and high actually buy



LOW: ANSWERING FROM THE DOORWAY

A sliver of scratchpad. Right answer for known shapes: transforms, formatting, lookups, execution work. On frontier models, low is dramatically better than its reputation.

MEDIUM: THE NOTEPAD

A real scratchpad. It sketches a couple of paths, prunes one, answers. The everyday driver for most real work, and OpenAI's default for a reason.


HIGH: THE WHITEBOARD SESSION

Plans, acts, verifies. On Claude, high is literally identical to not setting the dial at all, that's the factory setting. If you never touched the knob, this is what you've been paying for.

WHY THINKING COMPOUNDS

Models re-read everything they've written on every new token. Longer thinking makes every later token more expensive to produce. That's the snowball, and it's why the meter accelerates at the top rungs.

Extra high and max are specialist tools



XHIGH: THE WAR ROOM

Built for long-running, unattended agentic work. Anthropic's guidance pins it to tasks over 30 minutes. Powerful there, wasted on a five-second question.

MAX: THE COMMITTEE THAT WON'T ADJOURN

More churn, more branches revisited, and frequently the same answer high would have given. You're paying everyone in the room whether or not they add anything.


THE RECEIPTS FROM THE FIELD

One community ladder test: the same coding prompt finished in 12 minutes on low and roughly two hours on max, both producing working code. OpenRouter found capping effort cut costs 50 to 75 percent with accuracy inside the noise.

A FRESH-MODEL BONUS

Anthropic notes lower effort on its newest flagship often exceeds extra high on the previous generation. The bigger brain on a five-minute timer beats the smaller brain locked in the room all day.

Four ways the dial bites quietly



SILENT DOWNGRADES

Ask for a level a model doesn't support and some tools quietly run a lower one. No error, no warning. If a level matters, verify the tool honors it.

EFFORT IS INHERITED

In agent setups, subagents spin up at the parent's effort. One max decision at the top multiplies across every agent it spawns. That's how a monthly quota dies in an afternoon.


MODES ARE NOT EFFORT

Pro modes and ultra modes aren't higher rungs, they're more compute or literally multiple agents in a trench coat. Different product, different bill.

THE ENCRYPTED RECEIPT

On some models the thinking lands in your logs with the content stripped and an encrypted signature in its place. You're billed for tokens you can never read. Control how many there are.

Two seconds per task, any provider



PICK THE MODEL FIRST

Cheapest model that reliably finishes the job. Workhorse models for execution, frontier models for judgment. The dial is the second decision, never the first.

THEN START LOW

Known shape? Low. Real reasoning? Medium. Genuine ambiguity or architecture? High. Long unattended agentic run? Extra high. Max is for the rare linchpin where being wrong is the expensive outcome.

CLIMB ON EVIDENCE

Wrong or shallow output is the signal to move up one rung and rerun. That is the entire workflow. You'll land lower than you expect most days.

THE PAYOFF

This is how you stop paying for thinking your tasks never needed, without ever risking quality on the ones that do.

Every rung, decoded

The decision tree in slow motion. What each level is actually for, when to skip it, and the tell that it's time to climb.

LOW · ANSWER FROM THE DOORWAY

Use for: known shapes. Reformatting, extraction, file conversions, boilerplate, template filling, executing a plan that already exists. **Skip when:** the task has real multi-step logic. **Climb when:** steps go missing or the output contradicts itself. On the newest flagships, low regularly beats last generation's extra high.

MEDIUM · THE NOTEPAD

Use for: everyday coding, structured drafting, summaries that need judgment, moderate debugging. OpenAI ships here by default, and most real work lives here. **Skip when:** a pure formatting job could run on low. **Climb when:** it hand-waves through the hard part of the problem instead of working it.

HIGH · THE WHITEBOARD

Use for: planning small builds, refactors with real tradeoffs, final verification passes, genuinely ambiguous questions. On Claude, unset means high, this is the factory setting. **Skip when:** the task is routine, you're paying for deliberation it won't use. **Climb when:** the job is long, multi-file, and runs without you.

XHIGH · THE WAR ROOM

Use for: long-running agentic and coding work, the vendor guidance pins it at tasks over 30 minutes. Deep architecture, complex multi-file refactors, unattended runs. **Skip when:** you are sitting there waiting for the answer. **Watch out:** subagents inherit it, one xhigh parent can quietly spawn an xhigh fleet.


MAX · THE COMMITTEE

Use for: the rare linchpin where being wrong is the expensive outcome. System design, schema decisions, security-sensitive logic. **Skip when:** high already answers it, max frequently returns the same answer at a multiple of the cost, and the vendor's own docs warn it can overthink. **Rule:** if you can't say why this task deserves max, it doesn't.

THE FLOOR RUNGS, WHERE THEY EXIST

OpenAI's none and Gemini's minimal skip reasoning entirely: instant classification, tagging, and template filling at the lowest latency and price. Not every model has them, and asking for a missing rung can silently run you at a different level, so verify what actually executed. When speed is the product, the floor is a feature, not a compromise.

The labels don't survive model upgrades



SAME WORD, NEW MEANING

Anthropic's own docs: the new generation's medium roughly equals the old generation's HIGH. A full notch of drift, in writing, behind the same faceplate.

THE SPEND MOVES TOO

Independent benchmarking clocked the same max label spending differently across model generations. The dial position stayed, the bill changed.

IT'S OFFICIALLY A SIGNAL

Effort is a behavioral signal the model interprets, and vendors retune that interpretation every generation. OpenAI even documents that supported values are model-dependent.

THE RITUAL

New model, reset dial. Re-run your standard tasks at low and medium before assuming last month's setting still means what it meant. Takes ten minutes, saves a month of overpaying.

Same prompt, every rung. Let the receipts talk.

These are the exact control prompts from the video. The method matters as much as the prompts: the prompt is the control group, the dial is the only variable.

THE METHOD

One fresh session per run, never reuse context across levels. Make one folder per run (claude_low, codex_high) and launch from inside it. Paste the prompt verbatim, not one word added.

WHAT TO RECORD

Three numbers per run: wall-clock time, tokens or cost (Claude Code's /cost, or your tool's usage readout), and the score from the prompt's own built-in checklist. Then open the outputs side by side.

SETTING THE DIAL

Claude Code: `claude --effort low` (low, medium, high, xhigh, max). Codex CLI: `codex -c model_reasoning_effort=low`. Every provider's API takes the equivalent parameter.

WHAT YOU'LL SEE

All three prompts produce visual, self-scoring outputs, so the differences between rungs are obvious in seconds: missed data traps, escaped physics balls, broken dashboards, and wildly different bills.

THE THREE CONTROL PROMPTS, IN THIS DOWNLOAD

Prompt 1 builds a live X-sentiment dashboard with an AI scoring step and a synthetic-data fallback, so it works on any machine. Prompt 2 is a sales dashboard with planted data-quality traps, low effort misses them, high effort catches them. Prompt 3 is the classic rotating-hexagon physics stress test. Each prompt ships as its own copy-paste .txt file right in this download, paste one verbatim and let the rungs compete.

If you only keep one page, keep this one

EFFORT LEVELS AT A GLANCE (WORKS ON ANY PROVIDER)

LOW	MEDIUM	HIGH	XHIGH	MAX
known shapes: transforms, formatting, execution work	real reasoning, the everyday driver	planning, verification, genuine ambiguity	long unattended agentic runs, 30 minutes plus	the rare linchpin decision only

THE WORKFLOW

Pick the cheapest model that reliably finishes. Start low. Climb one rung only when the output tells you to. Verify anything that matters.

THE RITUAL

New model, reset dial. Labels do not carry over between generations, re-test your standard tasks at low and medium on day one.

Maxing out thinking is not maxing out intelligence. Saving your time and your tokens is.

THE ONE-LINE CHEAT SHEET

Model first. Start low. Climb on evidence. New model, reset dial. And never send the committee to answer a multiple-choice question.

GO DEEPER

Want to build systems like this yourself?

Inside the Early AI Dopters community, you get the full **LIVING** Claude Code course, updated every single week:

THE FUNDAMENTALS

Operator fundamentals (effort levels, subagents, dynamic workflows, plan handoffs) taught from zero.

THE EXACT PATTERNS

The decoding workflow from this guide, applied to real builds, walked through step by step in dedicated modules.

DIRECT ACCESS

Me and my team of coaches when you hit a wall, on prompting, agentic systems, and monetizing AI.

THE WORKING GROUP

A community of builders shipping with these models without burning their credits.

Join here: <https://www.skool.com/earlyaidopters/about>