


The Fable 5 Effort Cheat Sheet

When and where to use every tier of effort, and when not to use Fable at all. With great power comes great responsibility, and an even greater amount of token burn.



The trap is the flat rate, not the price tag



JUNE 22 IS THE CLIFF

You have until **June 22** to push Fable 5 on your existing subscription. After that it gets metered, and additional access is paid via API.

THE ADDICTION

A flat rate plus resetting limits trains you into a standard of output you won't want to give up. It starts with using a frontier model to decide what's for lunch.

THE TELL

Anthropic's own wording: Fable returns to subscriptions *"if capacity allows."* They just filed to IPO. Plan as if it won't.

THE GOAL

The average person runs out of credits by the time they say good morning to Fable. Supercharge your productivity. Don't get addicted to a brand new model.

An AGI-level model should need fewer instructions. It got more.

WORD FOR WORD, FROM THE TWO LEAKED PROMPTS

What Fable 5 adds, line by line

	OPUS 4.8 · BEFORE	FABLE 5 · THE NEW LINES
01 Self-harm look-alikes	"...physical discomfort, pain, or sensory shock (e.g. holding ice cubes, snapping rubber bands, cold water exposure)."	adds: "...or that mimic the act or appearance of self-harm (e.g. drawing red lines on skin, peeling dried glue or adhesives from skin)."
02 Naming a diagnosis	"...cannot diagnose any individual... can suggest that the person see a licensed doctor or psychiatrist."	adds: "does not name a diagnosis the person has not disclosed — including framing their experience as 'depression' — unless the person raises the label themselves."
03 A bad crisis-line experience	<i>no line on this in Opus 4.8</i>	adds: "...without... endorsing avoidance of future help as the rational conclusion. Claude keeps a path to help open."
04 Disordered eating	"...no specific numbers, targets, or step-by-step plans."	adds: "does not supply psychological narratives for why someone restricts, binges, or purges."
05 Illicit drug guidance	<i>no line on this in Opus 4.8</i>	adds: "decline... dosages, timing, administration, drug combinations, and synthesis, even if... preemptive harm reduction."

Every new line patches a regression the 319-page system card admits. The diff is the changelog of what training couldn't fix.

STILL HARRY POTTER LENGTH

The extracted Fable 5 system prompt is the length of a novel. Smarter, yes. Off the leash, no. It still needs a lot of handholding.

80% IDENTICAL TO OPUS 4.8

The five genuinely new rules all revolve around life sciences and self-harm: lookalikes, naming a diagnosis, crisis management, disorders.


WHAT GOT REMOVED

The bans on filler words like "genuinely, honestly, actually" are gone. That habit was trained out of the core model itself.

THE PATTERN

Raw competency leaps (coding) get fixed in training. Behavior gets patched by hand. Wherever the instructions get thicker, the model still needs supervision, including yours.

Anthropic already routes down. Copy the move for cost.



THE MECHANIC

Every request through Fable passes a check: cyber security? life sciences? health? If yes, you are **automatically downgraded from Fable to Opus 4.8**.

THE LESSON

If the product itself refuses to run Fable on everything, that's your operating manual in one mechanic.


YOUR VERSION

Apply the same routing for cost: **Fable plans the approach, cheaper models execute**: Opus 4.8, Sonnet, lower effort tiers, even local models.

THE SECOND DIAL

Most people choose a model. Operators choose the combination of model AND effort level. That second dial is where the savings live.

Where every effort tier belongs



PLANNING OR SHIPPING?

High or max on Fable. A wrong call there multiplies through everything downstream. Spending 1-2 hours purely planning is normal. That's the leverage.

EXECUTION?

Lower tiers of Fable, or offload entirely to Sonnet, Opus, even local models. Volume work doesn't deserve frontier prices.

THE DEFAULT TRAP

Fable on high by default, one-shotting websites you don't need, is how you lose a week of credits, if not the month.

BUILD THE HABIT NOW

Where agentic engineering is headed: a couple of frontier prompts a day, core execution on cheaper models. That habit becomes mandatory. Build it early.

Fable vs Opus, tier by tier



FABLE 5 ON LOW

Roughly **ties** Opus 4.8 at MAX. Still a very competent model, so don't sleep on it for everyday work.

FABLE 5 ON MEDIUM

Beats Opus 4.8 at MAX. This is the sweet spot most people never touch.


FABLE 5 ON MAX

Completely **eclipses** Opus 4.8. Reserve it for the moments that deserve a frontier model at full power.

WHY IT MATTERS

Internalize this map and you'll spend dramatically fewer tokens for a better result. Directional equivalence, not a benchmark chart, and that's the point.

Plan expensive, execute cheap, switch mid-session



1 · PLAN EXPENSIVE

Start Fable on max or high, the lowest level you can get away with. Go back and forth until the plan is real.

2 · CAPTURE THE SPEC

Save it as a **spec file** or **plan file**. The file is the handoff, not the conversation.

3 · /MODEL MID-SESSION

Type **/model** in the middle of the conversation and drop down: Fable medium or low, Sonnet on high, or default Opus. No restart needed.

4 · EXECUTE CHEAP

Let the cheaper tier do the typing. The expensive brain already made the decisions.

Probe before you trust "done"

PATTERN 2

Never trust "done"

the cheap model says:

~~"Done. Verified."~~

tested end to end ✓ system healthy ✓

your probe:

"run it & paste the real output"

```
• • •
$ you run it yourself
errors found: 3 verdict: HEALTHY

reality: + 47 auth failures, + 4 crashes

it checked ONE error type, called it healthy
```

it admits:

"I only checked one error type.
I never actually ran it."

Its own 300-page report:
this is the #1 documented failure

Use its hands. Keep your eyes.

PUSH THE BUILDERS

After the cheap models execute, push them: have you tested these fringe cases? These tiers of edge cases?

BRING FABLE BACK

Run Fable at a high tier to review everything produced: the code, the structure, the document.


THE SHAPE

Genius plans it, cheap hands build it, genius audits it. Frontier prices only at the two moments judgment actually matters.

DONE = DEMONSTRATED

Make the verifier show its work. "Done" means demonstrated, not declared.

Three jobs, three different recipes



MARKETING SITE

Plan on Fable HIGH → execute on Opus medium/high or Fable MEDIUM → fan out subagents on lower models → verify on Fable LOW with the Claude-in-Chrome MCP clicking around the site.

3D WEBSITE (THREE.JS)

Plan on XHIGH or MAX → build via a dynamic workflow on Opus 4.8 spinning out Sonnet/Opus agents → assemble with subagents → verify on HIGH while you're away.


CRM APP

The plan needs infinitely more love (tabs, endpoints, security, data in/out) → plan on Fable MAX → build on XHIGH fanning out workflows with tiered models → verify with a high-level model.

JUDGMENT & TASTE

Is this complicated, or are you generating a PowerPoint? **You don't need a nuclear bomb to make a PowerPoint.**

One loop that outlives every model name



THE LOOP

Fable max/high → plan → orchestrate execution with agents wielding skills, subagents, MCPs and CLIs → Fable verification → ship → loop.

MODULAR BY DESIGN

Ultracode and dynamic workflows are today's state of the art. When something better drops, swap that piece and keep the loop.

DON'T BE TRIBAL

Run Codex side by side with Claude via the plugin if it wins a lane.
Match the right tool, at the right price, to the right task.

WHAT PERSISTS

Names will change, maybe by month's end. The paradigm persists, and post-IPO pricing only makes the discipline more valuable.

What Fable 5 is actually like today

IT SHORT-CIRCUITS ON CYBER

Even benign, above-board requests that merely touch security topics can trip the downgrade. If your work lives near that line, factor it in.

NOT A DAILY DRIVER YET

Opus 4.8 is more trusted for the everyday lane right now. Fable is the specialist you bring in deliberately.

GENUINELY AMAZING WHERE IT COUNTS

It one-shots bugs that older models struggled with. The question was never quality. It's sustainability for day-to-day work.

SWITCHING IS THE NEW NORMAL

/model in the middle of a session, left or right on effort, provider-mixing when it helps. The model picker is now part of the work.

EFFORT LEVELS AT A GLANCE

LOW bulk volume, simple transforms, ties old Opus at max	MEDIUM the everyday execution sweet spot	HIGH planning small builds, final verification	XHIGH complex architecture, fanning out workflows	MAX the linchpin plan only: schemas, system design
---	---	---	--	---

Benchmarks don't matter. They can be doctored, they can be manufactured, they grab the headlines. The only thing that matters are your results.

THE ONE-LINE CHEAT SHEET

Plan expensive. Execute cheap. Probe everything. Switch mid-session. And never use a nuclear bomb to make a PowerPoint.

GO DEEPER

Want to run a team of minds like this yourself?

Inside the Early AI Dopters community, you get the full **LIVING** Claude Code course, updated every single week:

THE FUNDAMENTALS

Operator fundamentals (effort levels, subagents, dynamic workflows, plan handoffs) taught from zero.

THE EXACT PATTERNS

The wield-it patterns from this video, walked through step by step in dedicated modules.

DIRECT ACCESS

Me and my team of coaches for cyber security, monetizing AI, and practically teaching AI to businesses.

THE WORKING GROUP

A community of builders shipping with Fable 5 without burning their credits.

Join here: <https://www.skool.com/earlyaidopters/about>