

THE COMPLETE RAG LADDER CHEAT SHEET



From "Just Paste It" to Self-Organizing Intelligence Networks

LEVEL 0: CONTEXT DUMP

"Just Paste Everything"

How it Works: Copy your entire knowledge base (up to 700 pages) and paste it directly into the AI prompt along with the question.

Setup: 0 minutes

Cost: \$56-350/month (100 queries @ 200K tokens)

Accuracy: 95-100%

Speed: 5-10 seconds

Context Limit: ~700 pages (1M tokens)

Best For:

- Company wikis under 500 pages
- Legal firms reviewing contracts
- Small business product docs
- When accuracy > cost

Reality Check: Works perfectly for 90% of SMBs. No databases, no complexity, just paste and go.

LEVEL 1: PROMPT CACHING

"Same Thing, 75% Cheaper"

How it Works: First query costs full price, AI caches the knowledge base for 5-60 minutes, subsequent queries only pay for new questions.

Setup: 10 minutes

Cost: \$10-50/month (75% savings)

Accuracy: 95-100% (same as Level 0)

Speed: 2-5 seconds

Cache Duration: 5-60 minutes

Best For:

- Customer service rushes
- Morning question batches
- Teams working simultaneously
- Any clustered activity pattern

Pro Tip: Batch your queries within cache windows for maximum savings.

LEVEL 2: SMART CHUNKS

"AI Librarian Picks Chapters"

How it Works: Divide docs into logical sections (HR, Finance, Tech). Cheap AI (GPT-4.1 mini) picks relevant section, expensive AI only reads that section.

Setup: 1-2 days

Cost: \$15-25/month

Accuracy: 90-95%

Saves: 60-80% on costs

Best For:

- Clearly organized knowledge bases
- Distinct departments/categories
- Documents with minimal overlap
- When topics don't cross boundaries

Risk: Might miss context spanning multiple sections.

LEVEL 3: TRADITIONAL RAG

"The Card Catalog System"

How it Works: Convert docs to embeddings (mathematical fingerprints), store in vector database, search for similar fingerprints, retrieve 5-10 relevant paragraphs.

Setup: 2-4 weeks

Cost: \$50-100/month infrastructure

Accuracy: 60-70% ⚠️

Speed: Milliseconds

Scale: Millions of pages

Best For:

- Massive document collections (1M+ tokens)
- Simple Q&A
- When speed > accuracy
- Need source citations

Fatal Flaw: Misses context, only finds fragments, "lost in the middle" problem.

LEVEL 4: HYBRID RAG

"Semantic + Keywords = Better"

How it Works: Two parallel searches - semantic (finds concepts: "car" → "vehicle") + keyword (finds exact: "ERROR_2847"). Combines results.

Setup: 3-4 weeks

Cost: \$75-150/month

Accuracy: 75-85% (18.5% better than traditional)

Maintenance: 2x complexity

Best For:

- Technical documentation
- Medical/legal (specific terms + concepts)
- E-commerce (SKUs + descriptions)
- Mixed content types

Trade-off: Better accuracy but double the infrastructure.

LEVEL 5: CONTEXTUAL RAG

"Grabs the Whole Section"

How it Works: Instead of retrieving one sentence, grabs 2-3 paragraphs before and after. Maintains logical flow.

Setup: 4-5 weeks

Cost: \$100-200/month

Accuracy: 80-88%

Token Usage: 3-5x traditional RAG

Best For:

- Instruction manuals
- Step-by-step procedures
- Legal contracts with dependencies
- Medical procedures

Example: Instead of "Press reset button", returns entire 5-step procedure including prerequisites.

LEVEL 6: AGENTIC RAG

"AI Research Assistant"

How it Works: AI agent controls search, tries multiple strategies, evaluates results, refines queries, can search 5-10 times per question.

Setup: 4-8 weeks

Cost: \$150-300/month + unpredictable

Accuracy: 85-92%

Risk: Can spiral costs (10x per query)

Best For:

- Complex troubleshooting
- "Why" questions
- Vague user queries
- Research requiring detective work

Warning: Agent might search endlessly if not controlled.

LEVEL 7: MULTI-INDEX AGENTIC RAG

"Team of Specialist Librarians"

How it Works: Separate indexes for different content types (technical, email, financial, legal). Orchestrator agent routes to appropriate specialists.

Setup: 6-10 weeks

Cost: \$300-500/month

Accuracy: 88-94%

Indexes: 3-10 specialized

Best For:

- Large enterprises
- Diverse data types
- Cross-functional queries
- Multiple departments/systems

Example: "What did John say about Q3 budget?" → searches email + financial + project indexes.

LEVEL 8: GRAPHRAG

"The Relationship Map"

How it Works: Stores information as network of relationships. Entities are nodes, relationships are edges. Enables multi-hop reasoning.

Setup: 6-12 weeks

Cost: \$200-500/month

Accuracy: 95% on relationships, 70-80% general

Update Complexity: Nightmare level

Best For:

- Supply chain analysis
- "Who knows whom"
- Impact analysis
- Customer 360 views

Fatal Flaw: Adding one document = updating hundreds of relationships.

LEVEL 9: AGENTIC GRAPH RAG

"Self-Organizing Intelligence Network"

How it Works: Multiple agents maintain the graph - Discovery (finds relationships), Validation (verifies), Healing (fixes), Traversal (navigates), Learning (identifies patterns).

Setup: 3-6 months

Cost: \$500-2000/month

Accuracy: 90-96%

Team Required: Dedicated AI ops

Best For:

- Fortune 500 operations
- Intelligence agencies
- Large-scale research
- When relationships > documents

Reality: One rogue agent can corrupt everything. Only implemented once in 3 years.

THE DECISION FRAMEWORK

START HERE (90% OF CASES):

1. Under 200K tokens? → Level 0 or 1
2. Organized sections? → Level 2
3. Everything else → Start simple, level up only when hitting walls

KEY INSIGHTS:

72% of enterprise RAG implementations fail

- Most fail from over-engineering, not under-engineering
- The best system is the simplest one that works
- Update complexity kills more systems than initial setup

THE PATH OF LEAST RESISTANCE:

1. Try context dumping first (Level 0)
2. Add caching if volume justifies (Level 1)
3. Only go complex when you have real evidence you need it
4. Most businesses never need beyond Level 2

COST REALITY (100 QUERIES/MONTH, 200K TOKENS):

- Level 0-1: \$10-56/month
- Level 2-3: \$15-100/month
- Level 4-6: \$75-300/month
- Level 7-9: \$300-2000/month



Want access to an unlimited wealth of AI knowledge?

The Early AI-dopters community will take you to the next level.

Join today!

<https://www.skool.com/earlyaidopters/about>